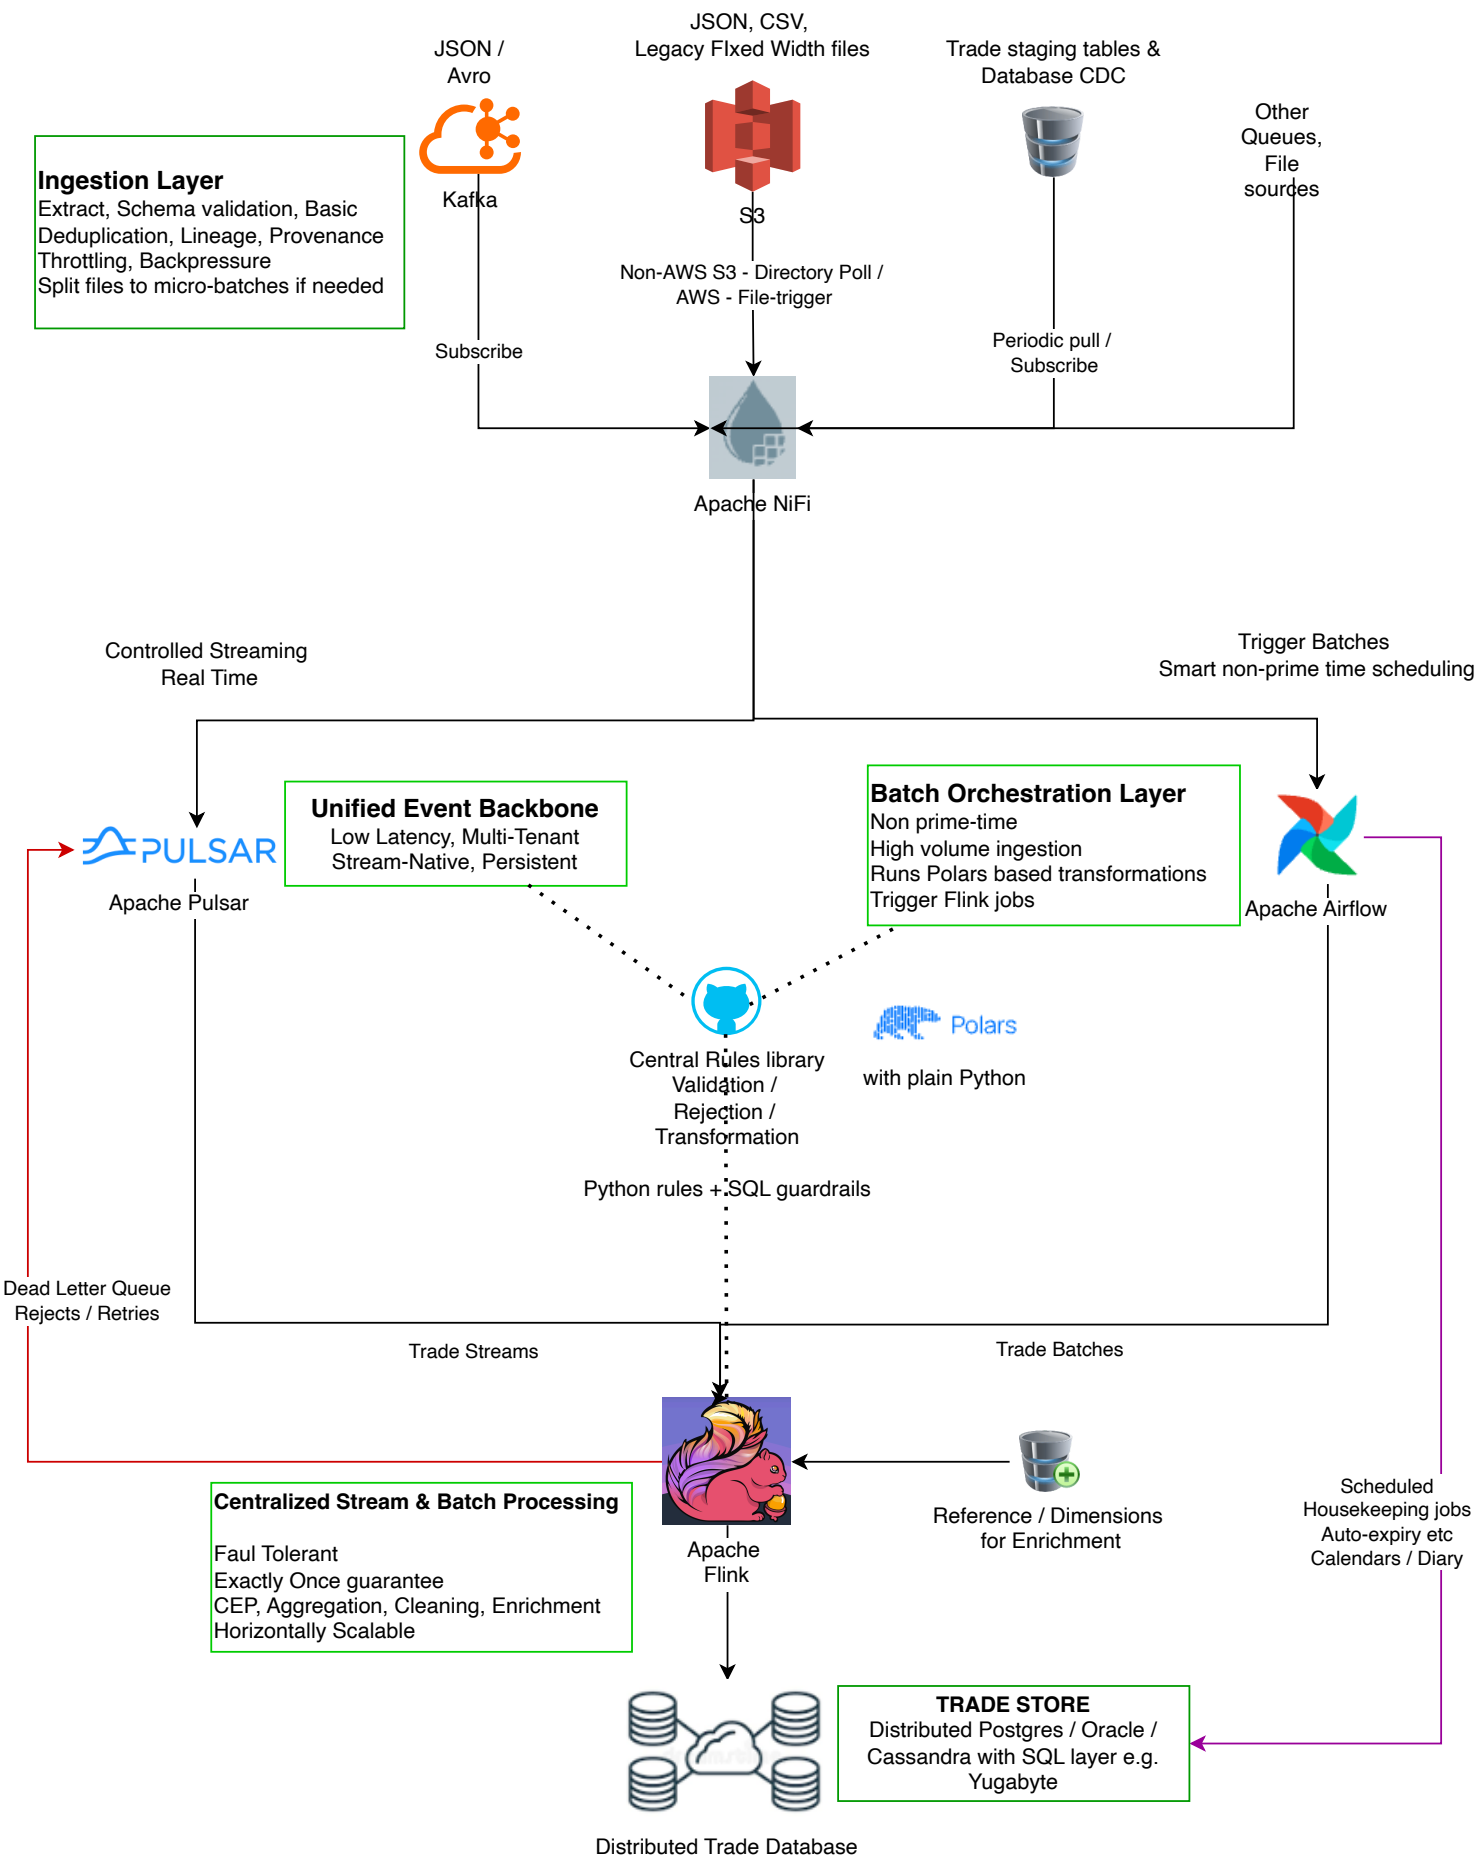


TRADE INGESTION PIPELINE



Flink Job Flow

- **Source:** defines a Flink Table `pulsar_trades` that consumes JSON from Pulsar (topic configurable). The table is converted to an append `DataStream` of Python rows.
- **Mapping:** each row is converted to a Python `dict` with keys `trade_id`, `version`, `counter_party_id`, `book_id`, `maturity_date`, `created_date`, `payload`.
- **Pre-filter (keyed state):** `TradePreFilterProcessFunction` is keyed by `trade_id` and keeps a `ValueState` named "`latest_active_version`":
 - **Rule 1:** If incoming `version` < `latest_active_version` → reject (side-output) with reason `RULE_1_LOWER_VERSION_THAN_ACTIVE`.
 - **Rule 2:** If incoming `version` == `latest_active_version` and `maturity_date` < today → reject with reason `RULE_2_MATURITY_BEFORE_TODAY_FOR_ACTIVE_VERSION`.
 - **Otherwise:** update stored latest version (if needed) and forward the trade to main stream.
- **DLQ side-output:** rejects from pre-filter are emitted to a side-output `rejects` and written to a Pulsar DLQ table `pulsar_dlg`.
- **Sink (DB Upsert):** `PostgresUpsertSink` uses `psycopg2` to `INSERT ... ON CONFLICT (trade_id, version) DO UPDATE ...`. If the DB raises an error containing `REJECTED_BY_DB_GUARDRAIL` (or variants), the sink will try to publish a structured DLQ message to the configured Pulsar DLQ (via `python pulsar client`). If that fails it logs an error and drops the record.
- **Fault-tolerance:** job enables Flink checkpointing; the keyed state (latest version) is persisted in checkpoints.